

Mastering Production RAG: Pipeline Order & Learning Roadmap

From Knowledge Preparation Fundamentals to Production-Grade Retrieval and Agentic Systems

Yes — your reasoning is correct, with one important clarification.

If your goal is to understand RAG deeply, from fundamentals to production-grade RAG, I recommend that you follow the pipeline order rather than jumping directly to Retrieval.

Your observation is actually very important:

Core Principle: Retrieval has high value, but you cannot fully understand Retrieval unless you understand what happens before it.

The RAG pipeline as a learning path

I would keep this order:

```
Stage 1 → Data Ingestion
↓
Stage 2 → Document Processing
↓
Stage 3 → Chunking
↓
Stage 4 → Embedding
↓
Stage 5 → Indexing
↓
Stage 6 → Retrieval ★
↓
Stage 7 → Reranking
↓
Stage 8 → Context Construction
↓
Stage 9 → Prompting
↓
Stage 10 → Generation
```

And then move into:

```
Stage 11 → Memory / Conversational RAG
Stage 12 → Evaluation
Stage 13 → Observability
Stage 14 → Security
Stage 15 → Production Architecture
```

Why I agree with your approach

Consider Stage 6 — Retrieval.


When you eventually learn:

"How does vector retrieval work?"

you should already understand:

- **What is being retrieved?** → Chunks
- **Where did those chunks come from?** → Document processing + chunking
- **Why can we search them semantically?** → Embeddings
- **Where are those embeddings stored?** → Vector index / vector database

So the complete mental model becomes:



That is a much stronger mental model than learning Retrieval in isolation.

But there is one important distinction

I would not say:

Retrieval is the "last step" of RAG.

It is the last step of the knowledge preparation/indexing side, but not the last step of the overall RAG pipeline.

Think of RAG as having two major phases.

```
Phase A – Prepare the knowledge
Documents
  ↓
Ingestion
  ↓
Processing
  ↓
Chunking
  ↓
Embedding
  ↓
Indexing
```

This happens before the user asks a question.


```
Phase B – Answer the user's question
User Query
  ↓
Retrieval
  ↓
Reranking
  ↓
Context Construction
  ↓
Prompt
  ↓
LLM Generation
  ↓
Answer
```

So Retrieval is actually the bridge between these two worlds.

That's why Retrieval is such a high-value topic.

Your learning strategy is therefore excellent

I would structure your journey like this:



Then:



What I recommend for you specifically

Since you want to go from basics → advanced → production, don't rush toward Retrieval just because it is technically interesting.

Instead:

Stage 1 — Data Ingestion

Learn:

- What ingestion means
- PDF ingestion
- Word
- HTML
- Markdown
- CSV
- JSON
- Database data
- APIs
- Cloud storage
- Batch vs streaming ingestion
- Production ingestion architecture

Then move to Stage 2.

Stage 2 — Document Processing

Learn:

- Parsing
- Cleaning
- Normalization
- Removing noise
- Metadata extraction
- Tables
- Headers/footers
- OCR
- Document structure

Then move to Stage 3.

Stage 3 — Chunking

This deserves deep study because poor chunking can destroy retrieval quality.

Learn:

- Fixed-size chunking
- Character-based
- Token-based
- Recursive chunking
- Sentence chunking
- Semantic chunking
- Document-aware chunking
- Parent-child chunking
- Hierarchical chunking
- Chunk overlap
- Chunk size optimization

Then move to Stage 4.

Stage 4 — Embeddings

Learn:

- What embeddings actually represent
- Dense vectors
- Dimensions
- Similarity
- Cosine similarity
- Embedding models
- Open-source embeddings
- Commercial embeddings
- Multilingual embeddings
- Domain-specific embeddings
- Embedding model selection
- Embedding evaluation

Then move to Stage 5.

Stage 5 — Indexing

Learn:

- What indexing means
- Vector indexes
- ANN
- HNSW
- IVF
- Vector databases
- Metadata indexes
- Keyword indexes
- Hybrid indexes
- Updates
- Deletes
- Persistence
- Scaling

Then you arrive at Stage 6 — Retrieval.

At this point, Retrieval will make much more sense.

And then go very deep into Retrieval

This is where I agree with you completely.

I would actually spend more time on Stage 6 than on most individual stages.

Your Retrieval journey could become:

Stage 6 – RETRIEVAL

- 6.1 Retrieval Fundamentals
- 6.2 Dense Retrieval
- 6.3 Similarity Search
- 6.4 Top-K Retrieval
- 6.5 Sparse Retrieval
- 6.6 BM25
- 6.7 Hybrid Retrieval
- 6.8 Metadata Filtering
- 6.9 Query Rewriting
- 6.10 Query Expansion
- 6.11 Multi-Query Retrieval
- 6.12 Query Decomposition
- 6.13 Query Routing
- 6.14 HyDE
- 6.15 Parent-Child Retrieval
- 6.16 Hierarchical Retrieval
- 6.17 Contextual Retrieval
- 6.18 Self-Query Retrieval
- 6.19 Graph Retrieval
- 6.20 Agentic Retrieval
- 6.21 Retrieval Evaluation

That is where your advanced RAG knowledge really starts to develop.

One more recommendation

Don't learn each stage as an isolated technology.

For every stage, keep asking:

"How does this decision affect the next stage?"

For example:


```
Bad Chunking
↓
Poor Embeddings
↓
Poor Index Quality
↓
Poor Retrieval
↓
Poor Context
↓
Poor LLM Answer
```

This is one of the most important concepts in production RAG.

Similarly:

```
Good Chunking
↓
Good Embeddings
↓
Good Index
↓
Good Retrieval
↓
Good Reranking
↓
Good Context
↓
Good Prompt
↓
Better Answer
```

So RAG quality is a pipeline problem, not merely an LLM problem.

My final recommendation

Follow the order. Don't jump to Stage 6 yet.

Your proposed learning strategy is the one I would recommend:

```
Stage 1 → Stage 2 → Stage 3 → Stage 4 → Stage 5 → Stage 6 → ...
```

But when you reach Stage 6 — Retrieval, don't treat it as just another chapter.

Treat it as a major specialization within RAG and go from:

basic retrieval → dense → sparse → hybrid → filtering → query transformation → advanced retrieval → reranking → graph → agentic retrieval → evaluation → production.

That will give you something much more valuable than knowing how to call a LangChain retriever:

You will understand why a RAG system retrieves the wrong information, how to diagnose the problem, and what architectural technique to use to fix it.

And that is exactly the kind of understanding you want for production-grade RAG and eventually Agentic AI.